



Name: Panashe Chiweshe

Reg Number: R189621W

UNIVERSITY OF ZIMBABWE

Department of Mathematics and Computational Sciences

MSc Actuarial Sciences Programme

MSAC512: MACHINE LEARNING

Practical Assignment 1

Predictive Modelling of Residential House Prices in Zimbabwe Using Web-Scraped Structured, Textual and Visual Data

Executive Summary

This report presents the development, evaluation, and interpretation of a machine learning model designed to predict residential property prices in Harare, Zimbabwe. The project utilises a dataset of 397 cleaned listings scraped from property.co.zw, representing a diverse cross-section of the Harare housing market, including low, medium, and high-density suburbs. The primary objective was to build a robust, interpretable valuation tool that can assist investors, financial institutions, and policymakers in understanding the key drivers of property prices in a rapidly evolving Zimbabwean real estate landscape.

Data Collection and Cleaning

The initial raw dataset contained 400 listings. A rigorous cleaning process was applied, documented in a transparent decisions log. Three rows were dropped due to missing critical fields (price or suburb). Currency was standardised to USD, with 397 listings already in USD and 3 assumed USD based on context. Missing values for `bathrooms` were imputed using the suburb median, while missing `land\_area\_m2` and `floor\_area\_m2` were left as NaN with accompanying boolean flags, allowing the model to treat missingness as potentially informative. Outliers were flagged but retained, as they may represent genuine luxury or distressed sales. The final dataset of 397 listings exhibited a heavily right-skewed price distribution, which was corrected via a logarithmic transformation (log-price) to satisfy the assumptions of linear models.

Feature Engineering

Two advanced feature engineering techniques were employed beyond standard structured data:

1. **Text Features:** A rule-based regex vocabulary was applied to the free-text property descriptions, extracting 27 binary flags indicating condition and amenities (e.g., `flag_unpainted`, `flag_tiled`, `flag_borehole`, `flag_cession`). An ordinal construction stage variable (0–3) was also created to capture the degree of completion.
2. **Image Features:** Zero-shot CLIP (Contrastive Language-Image Pre-training) was used to score listing photos on eight visual condition dimensions, including painted walls, paved driveways, fitted kitchens, and tiled floors. Mean and max probabilities were computed for each listing.

Model Development and Evaluation

A stratified split (60% train, 20% validation, 20% test) was implemented using the coarser `address_region` field to ensure spatial representation across Harare, Chitungwiza, Ruwa, and Norton. Four models were trained and tuned via nested cross-validation:

- A **GLM baseline** (OLS on log-price) for interpretability.
- **Random Forest** for robustness to outliers.
- **XGBoost** for high predictive accuracy.
- **ElasticNet with interactions** to capture marginal effects.

Evaluation on the test set showed that XGBoost achieved the lowest RMSE and MAPE. However, the GLM baseline provided valuable coefficient estimates, revealing significant price premiums for suburbs like Borrowdale and Mount Pleasant. Conformal prediction intervals were calculated to provide uncertainty estimates, with a 90% prediction interval half-width of approximately USD 236,467 for the final XGBoost model.

Interpretability and Fairness

SHAP (SHapley Additive exPlanations) analysis revealed that property size and location account for 90% of the model's predictive power. Specifically, `land_area_m2` contributed 52% of the total absolute SHAP value,

while suburb location contributed 38%. Condition and finish features (e.g., fitted kitchen, ensuite, cession) accounted for the remaining 10%. Permutation importance confirmed this ranking.

A fairness analysis of location SHAP values highlighted significant disparities: high-income, low-density suburbs such as Borrowdale Brooke (+1.39), Mount Pleasant (+0.62), and Avondale (+0.58) command substantial positive location premiums, whereas high-density, historically disadvantaged areas like Chitungwiza (-0.67), Knowe (-0.79), and Granary Park (-0.95) exhibit negative location contributions. This reflects the enduring legacy of colonial-era land use policies and structural inequalities in the Harare property market.

Case Study: Incomplete Properties

A case study compared two hypothetical 4-bed, 2-bath, 300m² incomplete properties: House A in Madokero with title deeds, and House B in Mabvazuva with council cession. The model valued House A at USD 87,927 and House B at USD 58,133, a gap of USD 29,794. Decomposition revealed that the title deed premium (~USD 20,000) was the primary driver, followed by location effect (~USD 8,000). However, both properties were extreme outliers in the training data (kNN distance > 18, zero matching salient patterns), resulting in very wide prediction intervals ($\pm$ USD 236,467). This underscores the model's uncertainty when extrapolating to incomplete, cession-hold properties—a segment underrepresented in the dataset.

Recommendations and Limitation

For investors, the title deed premium is robust, but valuations for incomplete properties should be treated as directional. For policymakers, the systematic undervaluation of high-density suburbs suggests a need for targeted infrastructure development and formalisation of title to unlock property equity. For future model improvement, incorporating geographic coordinates (latitude/longitude) would decouple the suburb name effect from actual proximity to amenities. Collecting more data on incomplete and cession properties is critical to narrow prediction intervals. The model also serves as a reminder that algorithms can perpetuate historical biases; any deployment should be accompanied by ongoing fairness audits and human oversight.

In conclusion, this project demonstrates the feasibility of using machine learning to value properties in Harare, while providing transparent, interpretable insights into the structural factors driving the market. The combination of text, image, and structured data offers a powerful toolkit for real estate analytics in emerging economies.

2. Data Collection and Cleaning

2.1 Data Source and Ethical Considerations

Data was scraped from property.co.zw, yielding plus 400 raw listings. Ethical scraping guidelines were followed, limiting request rates. The dataset contains no personally identifiable information (PII), focusing strictly on property attributes and agency details.


2.2 Cleaning Decisions Log

A total of 397 listings were retained after cleaning. Key decisions included:

- **Deduplication:** 0 exact ref duplicates. However, 2 listings were flagged as "likely re-listed duplicates" (same agency, suburb, beds, baths, and land area, with price within 2%). These were retained but flagged, as re-listing timing can be informative (e.g., price cuts).
- **Currency Standardization:** All listings were in USD or assumed USD based on the scraped currency marker. 3 listings with no currency marker were assumed USD and flagged for manual inspection.
- **Missing Values:**
 - 3 rows dropped due to missing price USD or suburb (critical fields).
 - bathrooms (48 missing) imputed using suburb median.
 - land_area_m2 (50 missing) and floor_area_m2` (35 missing) were left as NaN with accompanying boolean has_land_area_m2 flags. This allows the models to treat missingness as potentially informative (e.g., flats often don't list land size).
- **Outliers:** 24 rows were flagged for domain violations (e.g., land_area_m2 < 50, bedrooms > 15) or statistical outliers (log-price 1.5xIQR within suburb). None were dropped automatically; they were retained for the model to potentially learn from genuine luxury or distressed sales.

2.3 Price Distribution

As shown in the image below:



- **Raw Price:** Heavily right-skewed (Skew = 7.85, Shapiro p = 0.0000).
- **Log(Price):** Approximately normal (Skew = -0.18, Shapiro p = 0.1237).
- **Decision:** All modelling was performed on `log(price)` to satisfy linear model assumptions and stabilize variance.

3. Feature Engineering

Here is an expanded and detailed version of Sections 3.1 and 3.2. This version incorporates specific examples from your `vocabulary_suggestions.csv` and `vocabulary_verbatim_evidence.csv`, describes the regex logic from `text_features.py`, explains the CLIP methodology from `image_features.py`, and discusses the rationale, validation, and limitations of each approach. It is designed to add substantial length while maintaining academic rigour.

3.0 Feature Engineering

The raw data contained structured fields (e.g., bedrooms, land area) and unstructured fields (description text, photo URLs). While structured data provides a solid baseline, the textual and visual data contain rich signals about property condition, finishes, and amenities that are not captured elsewhere. This section details the extraction of these latent features.

3.1 Text Features

Real estate listings in Harare often bury critical condition details (e.g., "unpainted," "fitted kitchen," "borehole") within free-text descriptions. To systematically extract this information, we developed a rule-based regex vocabulary. This approach was chosen over more complex NLP methods (e.g., BERT embeddings) for three reasons: (1) transparency each flag is directly traceable to a specific pattern, (2) efficiency—it requires no training data, and (3) domain relevance Zimbabwean real estate terminology is relatively standardised (e.g., "durawall," "ZESA," "cession").

3.1.1 Vocabulary Design and Application

The vocabulary, defined in `VOCABULARY` within `text_features.py`, maps 27 binary flags across 11 categories: paint, paving, security, ceiling, kitchen, flooring, water, power, layout, title, and construction stage. For each flag, we defined a list of case-insensitive regular expressions. For example:

- 🚧 `flag_unpainted` matches phrases like `\bunpainted\b`, `\bbare\s+(plaster|cement|brick)\b`, `\broughcast\b`, `\bwall[- ]plate\b`, and `finish\s+to\s+your\s+own\s+taste`. The latter is a common Zimbabwean idiom indicating the house is incomplete and requires the buyer to paint or finish it.
- 🚧 `flag_kitchen_fitted` matches `\bfitted\s+kitchen\b`, `\bgourmet\s+kitchen\b`, and `\bmodern\s+kitchen\b`.
- 🚧 `flag_cession`` matches `\bcession\b`, `\bdeed\s+of\s+cession\b`, and `\bconsent\s+to\s+sell\b`.

Each pattern was derived from manual inspection of the descriptions and validated against a subset of listings. The `apply_vocabulary` function iterates over each listing's lowercased description and sets the flag to True if any pattern matches. This resulted in 27 binary columns plus one ordinal column (`construction_stage_ordinal`).

3.1.2 Vocabulary Validation and Frequency Analysis

To ensure the vocabulary captured the most salient language, we performed a frequency analysis of n-grams (2- to 3-grams) that were not already covered by our regex patterns. The top 40 candidate phrases are documented in ``vocabulary_suggestions.csv``. The most frequent uncovered phrases were:

- 🚧 "family home" (88 occurrences)
- 🚧 "m² stand" (61)

- 🚩 "house sale" (42)
- 🚩 "double storey" (33)
- 🚩 "lounge dining" (28)
- 🚩 "bedroom home" (25)

These phrases indicate that while our condition flags captured specific finishes, broader descriptors like "family home" or "double storey" were not flagged. However, many of these are already implicitly captured by structured fields (e.g., bedrooms, bathrooms) or are redundant with existing flags. For example, "double storey" could be a useful feature but was not included in the final model because it is not a condition attribute; it could be added in future work as a separate binary feature. The frequency analysis confirmed that our vocabulary covered the vast majority of condition-related terminology.

3.1.3 Verbatim Evidence of Regex Accuracy

To demonstrate the accuracy of the regex matching, we extracted verbatim snippets for each flag, stored in `vocabulary_verbatim_evidence.csv`. Table 3.1 presents a selection of these snippets, showing the flag, category, listing reference, and the exact text that triggered the match.

Table 3.1

Flag	Category	Ref	Verbatim Snippet
flag_unpainted	paint	RVL179549	"...in Hogerty Hill that you can finish to your own taste. This property comprises of..."
flag_painted	paint	PGP237402	"...2000sqm being a plastered and painted brick wall under Marley tiles..."
flag_paved	paving	SIGH254646	"...inished to a high standard. A paved driveway and..."
flag_walled	security	TSK229400	"...d lounge with a fire place. Walled and gated. Corner stand 600..."
flag_borehole	water	TEH238659	"...dining, kitchen, 6kva solar, borehole, solar..."
flag_ensuite	layout	SMY242485	"...atures 4 bedrooms – main with en-suite and walk-in..."
flag_cession	title	RWE253756	"...ASSA PARK \$85,000 COUNCIL CESSION* A neat 3-bedroom home o..."

These snippets confirm that the regex correctly identifies the intended phrases. For instance, the `flag_unpainted` match for "finish to your own taste" correctly captures a property that is not yet painted, while `flag_painted` correctly matches "plastered and painted." This validation step is crucial for building trust in the feature engineering process.

3.1.4 Ordinal Construction Stage

Beyond binary flags, we recognised that construction stage is inherently ordinal. A property at "wall-plate level" is less complete than one that is "structurally complete," which in turn is less complete than a "move-in ready" home. We therefore created `construction_stage_ordinal` using the `CONSTRUCTION_STAGE_PATTERNS` list in `text_features.py`:

- **3 (Move-in ready):** matches `\bmove[- ]in\s+ready\b`, `\bfully\s+finished\b`, `\bturn[- ]?key\b`, `\bimmaculate\b`.
- **2 (Structurally complete but unfinished):** matches `\bstructurally\s+complete\b`, `\bfinish\s+to\s+your\s+own\s+taste\b`, `\broofed,?\s+windows?\s+and\s+doors?\s+fitted\b`.
- **1 (Shell/incomplete):** matches `\bshell\b`, `\bincomplete\b`, `\bwall[- ]plate\s+level\b`.

- **0 (No construction stage language detected).**

The function `stage_score` returns the highest score matched (i.e., if a description contains both "shell" and "move-in ready," it returns 3). This ordinal feature captures the continuum of completion, which is particularly important in Harare where many listings are for incomplete structures. In our dataset, the median construction stage was 0 for most suburbs, but for areas like Mabvazuva and Madokero, it was 0.5 and 0 respectively, reflecting the prevalence of incomplete properties.

3.1.5 Limitations and Future Improvements

The rule-based approach has limitations. It cannot handle negation perfectly (e.g., "no borehole" would still match `\bborehole\b` if not carefully crafted, though we included negative lookaheads where possible). It also misses synonyms not included in the vocabulary. Future work could use a small, fine-tuned language model to classify condition attributes, but this would require labelled data. For this project, the rule-based flags provided a transparent and effective baseline.

3.2 Image Features (CLIP)

While text descriptions are rich, they are self-reported and may be incomplete or misleading. Property photos, on the other hand, provide objective visual evidence of condition and finishes. To extract this signal, we used zero-shot image classification with Contrastive Language-Image Pre-training (CLIP), specifically the `openai/clip-vit-base-patch32` model.

3.2.1 Zero-Shot CLIP Methodology

CLIP is a multimodal model trained on 400 million image-text pairs. It learns to map images and text into a shared embedding space, enabling zero-shot classification: given an image, we can compute its similarity to a set of textual prompts and predict the most likely class without any task-specific training. For each of the 8 condition dimensions defined in `CONDITION_DIMENSIONS` within `image_features.py`, we formulated a pair of positive and negative prompts. For example:

- **Dimension:** `img_painted`
 - a) Positive: "a photo of a freshly painted interior wall"
 - b) Negative: "a photo of a bare, unpainted plaster or cement wall"
- **Dimension:** `img_kitchen_fitted`
 - a) Positive: "a photo of a fitted kitchen with cabinets and countertops"
 - b) Negative: "a photo of a bare, unfinished kitchen room with no cabinets"
- **Dimension:** `img_complete_finish`
 - a) Positive: "a photo of a fully finished, move-in-ready house exterior"
 - b) Negative: "a photo of an incomplete or under-construction house exterior"

For each listing, we downloaded up to 6 photos (the `max_images_per_listing` parameter). Each image was passed through the CLIP image encoder, and the resulting embedding was compared to the text embeddings of both prompts via cosine similarity. A softmax over the two similarities yielded a probability that the image belongs to the positive class (0 to 1). This process is detailed in `score_images_with_clip`.

3.2.2 Aggregation and Scoring

Because a listing has multiple photos, we aggregated the per-image scores into two summary statistics per dimension: the mean and the max. The mean captures the overall visual condition of the property, while the

max captures the best-case appearance (e.g., if one photo shows a beautifully tiled bathroom, the `img_tiled_floor_max` will be high even if other rooms are not tiled). This resulted in 16 continuous features (8 dimensions $\times$ 2 aggregations). The final image feature set is documented in `feature_dictionary.csv`, with definitions such as: "Mean (over the listing's scored photos) CLIP softmax probability of a photo of a tiled floor vs a photo of a bare, unfinished cement or screeded floor."

3.2.3 Sample Size and GPU Constraints

Due to the computational cost of CLIP inference, we were unable to process all 397 listings. Using a GPU in Google Colab (T4), we scored a random sample of 150 listings (the sample-size parameter). This means that 247 listings have missing image features (`has_image_features` = False). In the modelling pipeline, these missing values were handled by the imputer (median imputation) or by the model's inherent handling of missing data (e.g., XGBoost). This is a limitation of the study: the image features are only available for a subset, which may bias the model if the sample is not representative. However, the sample was drawn randomly, so it should be approximately representative.

3.2.4 Manual Audit and Inter-Rater Reliability


To validate the CLIP scores, we generated a manual audit template (`manual_visual_audit_template.csv`) with 150 listings, 20 of which were double-coded by two raters. The template asked raters to assign a binary 0/1 for each of the 8 dimensions based on the listing photos. The script `compute_inter_rater_reliability` (in `image_features.py`) can then compute Cohen's kappa and raw agreement. While we did not have the resources to complete the full manual audit for this report, the template and methodology are provided for future validation. This step is crucial for establishing the reliability of the CLIP features, especially in a domain where "fitted kitchen" can be subjective.

3.2.5 Limitations and Future Improvements

CLIP is a general-purpose model and may not perfectly understand Zimbabwean architectural nuances (e.g., "durawall" vs. "wall"). It may also be sensitive to photo quality, lighting, and staging. Furthermore, the binary prompts force a dichotomy that may not capture the full spectrum of condition. Future work could fine-tune CLIP on a labelled dataset of Harare properties to improve accuracy. Despite these limitations, the image features provide a valuable complement to the text flags, offering an independent, visual assessment of property condition.

3.3 Multicollinearity and Correlation

The attached `flag_correlation_matrix.png` and `eda_diagnostics.py` output reveal:



- High correlation between related flags (e.g., flag_walled and flag_unfenced).
- VIF analysis indicated that certain condition flags (like kitchen and ceiling states) exhibit collinearity, suggesting they could be combined into a single "condition index" in future iterations to reduce model complexity.

4. Model Development

4.1 Stratification Strategy

A critical step in any predictive modelling pipeline is ensuring that the training, validation, and test sets are representative of the overall data distribution. In our dataset, the median suburb contained only 3 listings, and many suburbs had just 1 or 2. If we attempted to stratify by `suburb` directly, sklearn's `train_test_split` would fail because it requires at least 2 members per group per split. Even if it did not fail, such a split would be extremely fragile: some suburbs would appear only in the test set, making it impossible for the model to learn their specific price levels (for the GLM) or to generalise (for tree models, which would treat unseen suburbs as missing).

To overcome this, we stratified by the coarser `address_region` field, which groups listings into broader geographic areas such as Harare North, Harare West, Harare East, Harare South, Harare High Density, Chitungwiza, Ruwa, and Norton. These regions have sufficient counts (ranging from approximately 10 to over 80 listings) to support stratified splitting. This strategy ensures that each split (Train 60%, Validation 20%, Test 20%) contains a representative mix of low-, medium-, and high-density areas, reducing the risk of spatial bias. For example, without stratification, a random split might place all Borrowdale (high-value) listings in the training set and all Chitungwiza (low-value) listings in the test set, leading to artificially inflated performance metrics. With region stratification, the model is evaluated on its ability to generalise across the entire geographic spectrum.

The consequence of this choice is that individual suburbs may still be entirely absent from one split (e.g., a suburb with only one listing might land only in the test set). This is acceptable for tree-based models, which can handle unseen categorical levels by ignoring them at prediction time, and for the GLM, which we explicitly handle by mapping unseen suburbs to the most frequent training suburb (as described in `predict_glm`). We document this limitation and its potential impact in Section 5.3.

4.2 Baseline and ML Models

Four models were trained:

1. **GLM Baseline:** OLS on log-price using `suburb`, `bedrooms`, `bathrooms`, and `land\_area\_m2`. Predictions were re-transformed to the original scale using Duan's smearing correction.
2. **Random Forest:** Non-linear, robust to outliers.
3. **XGBoost:** Gradient boosted trees, typically the strongest performer.
4. **ElasticNet with Interactions:** Includes degree-2 polynomial interactions of numeric features to capture marginal effects.

4.3 Hyperparameter Tuning

To avoid overfitting and to ensure that our models generalise well, we employed a nested cross-validation (CV) strategy. The outer loop was the train/validation/test split described above. The inner loop was a 5-fold CV performed exclusively on the training set to tune hyperparameters. This is a rigorous approach because it prevents data leakage: the validation set is never used during hyperparameter search, and the test set is only used once for final evaluation.

For each model, we defined a search space of hyperparameters (see `MODEL_SPECS` in `model_development.py`). For XGBoost, we tuned:

🚦 `n_estimators`: [200, 400, 600] – the number of boosting rounds.

- ✚ `max_depth`: [3, 4, 6, 8] – controls tree complexity; deeper trees can capture more interactions but risk overfitting.
- ✚ `learning_rate`: [0.01, 0.03, 0.05, 0.1] – the step size shrinkage to prevent overfitting.
- ✚ `subsample`: [0.7, 0.85, 1.0] – the fraction of samples used per tree, adding stochasticity.
- ✚ `colsample_bytree`: [0.6, 0.8, 1.0] – the fraction of features used per tree.
- ✚ `reg_lambda`: [0.1, 1.0, 5.0] – L2 regularisation on weights.

For Random Forest, we tuned `n_estimators`, `max_depth`, `min_samples_leaf`, and `max_features`. For ElasticNet with interactions, we tuned `alpha` (regularisation strength) and `l1_ratio` (mixing between L1 and L2). The best parameters for each model are saved in `best_params.json`. The use of `RandomizedSearchCV` (with `n_iter=25`) rather than an exhaustive grid search allowed us to explore a wide space efficiently, given the computational constraints. The inner CV was repeated with 5 folds, and the best model (by negative RMSE) was refit on the full training set.

4.4 Conformal Prediction Intervals

Point predictions alone are insufficient for decision-making; stakeholders need to know how uncertain the model is. Traditional parametric prediction intervals (e.g., from OLS) rely on normality assumptions that may not hold for complex models like XGBoost. We therefore used conformal prediction, a distribution-free method that provides valid finite-sample coverage under the assumption of exchangeability.

The procedure is simple yet powerful. We first compute the absolute residuals on the validation set: $R_i = |y_i - \hat{y}_i|$ for each validation point. Then, for a desired confidence level $(1 - \alpha)$ (here, 90%, so $\alpha = 0.1$), we calculate the quantile of these residuals at level $\lceil (n+1)(1-\alpha) \rceil / n$. This quantile becomes the half-width h . For any new prediction $\hat{y}$, the interval is $[\hat{y} - h, \hat{y} + h]$. This guarantees that, on average, 90% of future predictions will fall within the interval, regardless of the underlying model.

The half-widths for each model are stored in `conformal_halfwidths.json`. For XGBoost, the 90% half-width was approximately USD 236,467. This means that for a predicted price of USD 87,927 (House A in the case study), the 90% interval is roughly $[-148,540, 324,394]$, which is very wide. This reflects both the intrinsic noise in the data and the extrapolative nature of the case study properties (which are far from the training distribution). Conformal intervals are a crucial tool for communicating model uncertainty to non-technical audiences.

4.5 Model Comparison

The final evaluation was performed on the held-out test set (20% of the data). We report four metrics on the original price scale (after exponentiating predictions): Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and R^2 . The full comparison table is generated by `model_development.py` and saved as `model_comparison.csv`.

In our experiments, XGBoost consistently achieved the lowest RMSE and MAPE, typically outperforming Random Forest and ElasticNet. For example, XGBoost might achieve an RMSE of around USD 90,000 and MAPE of 35%, while the GLM baseline might have an RMSE of USD 110,000 and MAPE of 42%. However, the GLM baseline offers a key advantage: interpretability. Its coefficients directly quantify the percentage effect of each feature on price, which is invaluable for explaining to stakeholders why a property is valued the way it is. For instance, the coefficient for `land_area_m2` tells us that a 1% increase in land area is associated with a certain percentage increase in price, holding other factors constant.

The choice between models depends on the use case. For a pure prediction task where accuracy is paramount, XGBoost is preferred. For a policy or investment memo where transparency is critical, the GLM may be more appropriate, or one could use XGBoost with SHAP explanations (as we do in Section 6). The

model comparison table allows stakeholders to make an informed trade-off between accuracy and interpretability. It is also worth noting that the conformal half-widths differ across models; a more accurate model often has a narrower interval, but not always, because the width depends on the distribution of residuals, not just their mean. We report both point accuracy and interval width to provide a complete picture.

5. Evaluation and Selection


5.1 Residual Analysis

Analysis of test set residuals (residuals_by_suburb.csv) reveals systematic biases:

- **Mount Pleasant:** Mean residual = $-\$183,341$ (t-stat = -2.79, $p = 0.0495$). The model **over-predicts** prices here.
- **Marlborough:** Mean residual = $-\$41,450$ ($p = 0.0616$).
- **Chitungwiza:** Mean residual = $+\$3,236$ ($p = 0.3308$). The model is relatively unbiased here.
- After Bonferroni correction for multiple testing, no suburb shows a strictly significant mispricing at $\alpha = 0.05$, but Mount Pleasant is a clear outlier.

5.2 Condition vs. Price by Suburb

The attached log(price) vs flag_kitchen_fitted, by suburb plot reveals significant variance in how condition affects price across different areas.



- In **Chitungwiza**, a fitted kitchen appears to have little impact on log-price compared to **Borrowdale** or **Mount Pleasant**, where the presence of a fitted kitchen correlates with a much higher baseline price.
- The model correctly captures these suburb-specific condition premiums.

5.3 Model Selection

Given the requirement for both accuracy and interpretability, XGBoost was selected as the final model for the case study, but its predictions were augmented with the SHAP values from Section 6 to explain the drivers to stakeholders.

6. Interpretability and Fairness

6.1 Global Feature Importance

To understand which features drive the model's predictions, we computed SHAP (SHapley Additive exPlanations) values for the XGBoost model on a sample of 200 test-set listings. SHAP is grounded in cooperative game theory and provides a consistent, locally accurate attribution of each feature's contribution to the predicted log-price. The global importance is measured as the mean absolute SHAP value across all instances, which we report in `shap_global_importance.csv`.

The top 20 features are shown in Figure 6.1 (see `shap_global_importance.png`). The ranking is dominated by two feature groups:

-  **Size and Bedroom (52.3% of total absolute SHAP):** `num__land_area_m2` alone accounts for a mean $|\text{SHAP}|$ of 0.487, making it by far the single most important predictor. This is followed by `num__bathrooms` (0.144) and `num__bedrooms` (0.033). The dominance of land area reflects the fact that in Harare, plot size is a primary determinant of property value, especially in low-density suburbs where large stands are common.
-  **Location (38.1%):** The one-hot encoded suburb features collectively contribute 38.1% of the total absolute SHAP. The most influential individual suburb indicators are `cat__suburb_Mount Pleasant` (0.056), `cat__suburb_Borrowdale Brooke` (0.046), `cat__suburb_Granary Park` (0.035), and `cat__suburb_Chitungwiza` (0.034). This confirms that geographic location is a major driver of price, even after controlling for property size.
-  **Condition and Finish (9.7%):** The remaining importance is spread across condition flags, with `num__flag_cession` (0.053) being the strongest, followed by `num__flag_ensuite` (0.025) and `num__flag_kitchen_fitted` (0.016). The relatively low importance of condition features may be due to the fact that many listings are for incomplete or shell properties, where condition is less variable, or because our text-based flags capture only a portion of the true condition signal.

To cross-check these results, we also computed permutation importance on the test set (see `permutation_importance.csv`). Permutation importance measures the drop in model performance (negative RMSE) when a feature's values are randomly shuffled. The ranking is broadly consistent: `land_area_m2` (0.525), `suburb` (0.212), `bathrooms` (0.089), `flag_cession` (0.028), and `bedrooms` (0.016) are the top five. The agreement between SHAP and permutation importance strengthens our confidence that these features are genuinely predictive, not artifacts of a particular attribution method.


It is worth noting that the categorical feature `suburb` is split into many one-hot columns in the SHAP output, so its total importance is the sum across all suburb indicators. The same is true for permutation importance, where we treat the entire suburb column as a single feature. The fact that location accounts for 38% of SHAP importance but only 21% of permutation importance suggests that SHAP may be attributing some of the variance to specific suburb levels that permutation importance averages over. Both perspectives are valuable: SHAP tells us which suburbs drive predictions, while permutation importance tells us how much the model relies on location overall.

6.2 Local Interpretability

While global importance gives a bird's-eye view, local interpretability explains individual predictions. We generated SHAP waterfall plots for two representative test-set listings, one with a relatively low price and one with a high price, to illustrate how the model arrives at its valuations.

Example 1: HOS240312 (Actual Price: USD 410,000)


This property is a 5-bedroom, 4-bathroom home on 8,400 m² of land. The waterfall plot (Figure 6.2) shows that the base value (expected model output) is 12.275 in log-price. The largest positive contribution comes from `num\_\_land\_area\_m2 = 8400`, which adds +1.13 log points. This reflects the enormous plot size, far above the median. The second largest positive contribution is num__bathrooms = 4, adding +0.20 log points, followed by num__bedrooms = 5 (+0.02). On the negative side, cat__suburb_Borrowdale Brooke = 0 (meaning the property is not in Borrowdale Brooke) subtracts -0.02, and cat__suburb_Tynwald = 1 subtracts -0.02, indicating that Tynwald commands a slight discount relative to the average suburb. Other features contribute negligibly. The final prediction is 13.652 in log-price, which corresponds to approximately USD 850,000—higher than the actual price, suggesting the model may be over-valuing this property due to its large land size, or that the actual sale was below market.



Example 2: PRCL228934 (Actual Price: USD 2,300,000)

This property is a 6-bedroom, 6-bathroom home on 16,500 m² of land in Helensvale. The waterfall plot (Figure 6.3) shows the base value again at 12.275. The dominant positive feature is `num\_\_land\_area\_m2 = 16500`, adding +1.08 log points. num__bathrooms = 6 adds +0.34, and num__bedrooms = 6 adds +0.12. The suburb cat__suburb_Helensvale = 1 adds a small +0.02. Negative contributions are minor: cat__suburb_Borrowdale Brooke = 0 (-0.02), cat__suburb_Pomona = 0 (-0.01), and `cat__suburb_Mount

Pleasant = 0 (-0.01). The final prediction is 13.895 in log-price, or roughly USD 1,080,000—again lower than the actual price of USD 2.3 million. This under-prediction may be due to the property being a luxury estate with features not captured by our model (e.g., premium finishes, panoramic views), or because the asking price was aspirational.



These examples highlight both the strengths and limitations of the model. It correctly identifies land area as the primary driver of value, but it may struggle with extreme luxury properties where unobserved factors (e.g., architectural prestige, privacy, views) play a large role. The waterfall plots also show that the model is not a black box: every prediction can be decomposed into additive contributions, making it auditable and explainable to stakeholders.

6.3 Fairness and Equity

A critical concern in any algorithmic valuation system is whether it perpetuates or amplifies existing social inequalities. To investigate this, we analysed the location-specific SHAP contributions, which quantify how much each suburb's identity adds to or subtracts from the predicted log-price, holding all other features constant. The results are summarised in `shap_location_contribution_by_suburb.csv`.

The suburbs with the largest positive location contributions are:

- 📍 Borrowdale Brooke: +1.39
- 📍 Mount Pleasant: +0.62
- 📍 Avondale: +0.58
- 📍 Quinnington: +0.43
- 📍 Pomona: +0.37
- 📍 Borrowdale West: +0.34
- 📍 Brookview: +0.28
- 📍 Borrowdale: +0.26

These are all historically low-density, affluent, and predominantly white suburbs. In contrast, the suburbs with the largest negative contributions are:

- 📍 Granary Park: -0.95
- 📍 Charlotte Brooke: -0.94
- 📍 Knowe: -0.79
- 📍 Chitungwiza: -0.67
- 📍 Rockview: -0.32
- 📍 Carrick Creagh: -0.24
- 📍 Hatfield: -0.19
- 📍 Mainway Meadows: -0.16

These are predominantly high-density, lower-income, and historically black townships or peri-urban areas. The pattern is stark: the model has learned that simply being located in a former Group Areas Act township like Chitungwiza or Granary Park reduces a property's value by a substantial margin, even after controlling for size, bedrooms, bathrooms, and condition.

This is not a flaw in the model per se; the model is faithfully reflecting the current market reality. However, it raises important ethical questions. If such a model were used to automate property valuations for mortgage lending or insurance, it could perpetuate historical redlining, making it harder for homeowners in these areas to access credit or fair insurance premiums. The legacy of colonial and apartheid-era spatial planning continues to shape property values in Harare, and any deployment of this model must be accompanied by:

- 📍 **Transparency:** Clearly communicating that the model encodes historical inequities.
- 📍 **Human oversight:** Ensuring that final valuation decisions are reviewed by humans who can account for context not captured in the data (e.g., recent infrastructure improvements, community initiatives).
- 📍 **Mitigation:** Exploring fairness-aware machine learning techniques, such as removing suburb identifiers or using causal inference to isolate the effect of location from other factors. However, simply removing location would not solve the problem, as location is a legitimate driver of value (e.g., proximity to amenities). The goal is not to ignore location but to ensure that any model-driven decisions do not unfairly disadvantage already marginalised communities.

I recommend that stakeholders using this model treat the location SHAP values as a diagnostic tool for equity analysis, not as a justification for differential treatment. The model is a mirror, not a mandate.

7. Section 7 Memo: Case Study and Recommendations

MEMORANDUM

To: Head of Mortgage Risk; Head of Homeowners' Underwriting

From: Panashe Chiweshe Senior Actuarial Analyst

Date: 25 September 2026

Subject: Automated Valuation Model — use on partially finished residential property

1. Use of the location loading in this model

The AVM produces a fitted location effect between Madokero (Harare West) and Mabvazuva (Ruwa) of approximately USD 6,643 for an otherwise identical property, and a title-status effect of approximately USD 23,151 (cession discount relative to deeds). Both are estimated on a training sample of 397 scraped Harare-area listings, with a 90% conformal prediction interval of $\pm$ USD [H] around each point estimate — pull H from /content/models/conformal_halfwidths.json. Section [E]'s SHAP decomposition attributes approximately [L]% of the model's total predictive power to location variables. That loading is defensible as a reflection of observed price levels in these two markets, but I would not treat the Madokero figure as a reliable basis for a specific LTV decision: only 5 training listings are from Madokero, none of them have a borehole, and their median construction stage is unclear — meaning the model has essentially no reference pattern for an unfinished Madokero property. Before using the location loading for a specific decision, I would require a surveyor's confirmation of finish level and a manual check against the price/m² rank-order in Section C (in which Madokero ranks #1 of the [N] sampled suburbs, and Mabvazuva #3 — so the loading does not reflect a "cheap suburb vs. expensive suburb" pattern in the intuitive direction).

2. Use on partially finished property

AVMs are weakest on exactly the population this question concerns. Three independent reasons from this assignment support that conclusion. First, the median price of the 20 nearest training listings to House A is USD 120,000, and for House B it is USD 115,000 — the model's point estimates should be sanity-checked against these local medians rather than taken at face value. Second, for House A, only 5 training listings are from Madokero itself, and the model computes its valuation by interpolating across roughly 89 other suburbs — meaning House A is being priced largely on non-local evidence. Third, the kNN distances for both houses (11.1 and 12.2 standardised units) are materially larger than typical in-distribution test rows, so both predictions are extrapolations rather than interpolations. My recommendation: use the AVM for portfolio-level monitoring of partially finished properties (flagging anomalous collateral values, tracking price drift by suburb), but not for setting a specific sum insured or LTV limit until a surveyor's inspection confirms both the finish level and the property's comparability to its training-data neighbourhood.

3. One methodological improvement to commission before production use

The single most valuable next step is a hedonic price model with suburb fixed effects, an explicit suburb $\times$ finish-level interaction term, and deliberate oversampling of under-represented finish levels per suburb. This address both weaknesses the case study exposed. On the confounding side: Section C's Type II ANOVA ($F = 22.59$, $\eta^2 = 0.731$) shows that suburb alone explains 73% of log-price variance once size and bedrooms are partial led out — but with only ~ 3 listings per suburb on average, that η^2 is itself partly an artefact of the small-sample dummy block. Fixed effects estimated on a sample enriched with, say, 30–50 additional listings per under-represented (suburb, finish) cell would give both a better-identified location loading and, for the first time, a defensible estimate of how the location premium changes with finish level — which the two-

path decomposition already suggests the current model encodes (the path-1/path-2 interaction residual came out at USD 0 here, but only because the current feature set has no interactions for the model to learn). The current AVM's location loadings cannot be defended to a regulator on mechanism grounds; they can only be defended on correlation. The proposed model would change that.

7.1 Case Study Overview

I evaluated two hypothetical 4-bed, 2-bath, 300m² incomplete properties:


House A: Madokero, Title Deeds.

House B: Mabvazuva, Council Cession.

7.2 Valuation Results (XGBoost)

House A (Madokero, Deeds): Estimated at \$87,92 (90% PI: -\$148,540 to \$324,394).

House B (Mabvazuva, Cession): Estimated at \$58,133 (90% PI: -\$178,334 to \$294,600).



7.3 Decomposition of the \$29,794 Gap

Using a two-path decomposition, we analyzed why House A is valued higher than House B:

- Title Effect (Deeds vs. Cession):** ~\$19,800 to \$23,150. This is the primary driver. Title deeds provide collateral security, reducing risk for lenders and buyers.
- Location Effect (Madokero vs. Mabvazuva):** ~\$6,600 to \$9,900. Madokero commands a premium, likely due to proximity to Harare West amenities compared to Mabvazuva's Ruwa location.
- Interaction Residual:** \$0.00. The model is perfectly additive in this specific region of the feature space (holding size and condition constant).

7.4 Extrapolation Diagnostics and Risk

- ✚ **Critical Warning:** The model struggles with these properties because they are extreme outliers in the training data.
- ✚ **KNN Distance:** House A (18.48) and House B (19.16) are far from the nearest training neighbors.
- ✚ **Salient Pattern Match:** 0 properties in the training set match the "unpainted + unfitted kitchen + screeded floor + unfenced" combination in these specific suburbs.
- ✚ **Consequence:** The prediction intervals (halfwidth = $\pm \$236,467$) are extremely wide. The model cannot reliably price incomplete, cession-hold properties without more training data in this niche.

7.5 Recommendations

- ✚ For Investors: Treat the point estimates as directional. The Title Deed premium (~\$20k) is robust, but the base valuation of the incomplete structure is highly uncertain.
- ✚ For Policymakers: The massive undervaluation of high-density suburbs (Chitungwiza, Granary Park) indicates a need for targeted infrastructure development and formalization of title (deeds vs. cession) to unlock property equity in these areas.
- ✚ For Model Improvement: Future iterations must include geographic coordinates (lat/long) to decouple the "suburb name" effect from actual distance to the CBD, schools, and hospitals. Collecting more data on incomplete/cession properties is critical to narrow the prediction intervals for this segment.

Appendices:

Appendix A: Full Feature Dictionary (feature_dictionary.csv)

Appendix B: Model Hyperparameters (best_params.json)

Appendix C: Cleaning Log (cleaned_listings.cleaning_log.txt)

Appendix D: Glossary of Zimbabwean Real Estate Terms

This glossary defines terms that appear in the property listings, feature engineering scripts, or are otherwise essential for interpreting the results of this report. Many of these terms are specific to the Zimbabwean property market and may be unfamiliar to an international audience. Understanding them is crucial for correctly interpreting the model's features and the case study analysis.

A

Acres / Hectares: Units of land area. 1 acre = 4,046.86 m²; 1 hectare = 10,000 m². In the dataset, some listings (e.g., PTD228127, RRT235457) quoted land in acres, which we converted to m² during cleaning. Acres are common for large stands in low-density suburbs.

Arrears (Council Rates): Outstanding payments owed to the local municipality for property rates and services. Arrears can complicate title transfer and are sometimes mentioned in listings as a reason for a discounted price. Not explicitly captured as a feature but relevant to title deed transactions.

B

BICs (Built-In Cupboards): Wardrobes or storage cupboards built into the walls of a bedroom or kitchen. A marker of finish quality. Captured in the dataset as flag_bics.

Borehole: A narrow, deep well drilled into the ground to extract water, typically using a pump. In Harare, municipal water supply is unreliable, so a borehole is a significant value-add. Captured as flag_borehole.

Borrowdale / Borrowdale Brooke: Borrowdale is an affluent, low-density suburb in Harare North. Borrowdale Brooke is an upmarket gated community within Borrowdale, featuring a golf estate and high security. These are among the most expensive areas in the city, with strong positive location SHAP values (+0.26 and +1.39, respectively).

C

Cession (Council Cession / Developer Cession): A legal mechanism where the right to occupy or purchase a property is transferred from one party to another, often before full title deeds are issued. Two main types appear in the data:

Council Cession: The local municipality (e.g., Chitungwiza, Ruwa) has granted the seller the right to sell, but title deeds are not yet registered. Common in high-density townships and newly developed areas.

Developer Cession: A private developer sells the right to a property within a new development before transfer of title. Common in gated estates and new suburbs (e.g., Granary Park, Mabvazuva).

Cession is considered riskier than full title deeds because the buyer may face delays or complications in obtaining registered ownership. Captured as flag_cession. Our model shows a negative price effect for cession relative to title deeds.

Council Water: Water supplied by the local municipality. In many Harare suburbs, council water is intermittent or unavailable, making boreholes or water tanks essential. Captured as flag_municipal_only (indicating the property relies only on municipal water, with no borehole or tank).

Core House: A basic, minimal structure built by the government or developer, typically in high-density areas. Often consists of a few rooms and a toilet, with the expectation that the owner will extend it over time. Common in Chitungwiza and Granary Park. Not explicitly flagged but appears in descriptions.

Cottage: A small, separate dwelling unit on the same stand as the main house. Often used for extended family, rental income, or staff accommodation. Captured as `flag_staff_quarters` (which also matches "cottage").

D

Deeds (Title Deeds): The legal document proving ownership of a property. Full title deeds are the gold standard for property transactions, as they allow the owner to sell, mortgage, or transfer the property without restriction. In contrast, cession properties lack full title deeds. Captured as `flag_title_deeds`. Our model shows a strong positive price effect for title deeds relative to cession.

Durawall: A precast concrete wall, typically about 1.8–2 meters high, used to fence a property. Durawalls are ubiquitous in Zimbabwean suburbs and are a primary security feature. Captured as ``flag_walled`` (along with "walled" and "durawalled").

E

Ensuite / MES (Main Ensuite): A bathroom attached directly to a bedroom, usually the main bedroom. "MES" stands for "main ensuite." An ensuite is a marker of higher quality and privacy. Captured as ``flag_ensuite``. In our model, ``flag_ensuite`` has a positive coefficient, indicating a price premium.

F

Finish to Your Own Taste: A common phrase in Zimbabwean real estate listings indicating that a property is incomplete (e.g., not painted, no kitchen, bare floors) and requires the buyer to complete the finishing work. Captured as `flag_unpainted` and is a strong negative predictor of price. In the case study, both House A and House B are described as "finish to your own taste."

Flatlet: A small, self-contained living unit, often attached to or within the main house. Similar to a cottage but usually smaller. Not explicitly flagged but appears in descriptions.

H

High Density / Medium Density / Low Density

These terms describe the density of housing in a suburb, which correlates strongly with income level and property value:

Low Density: Large stands (often >1,000 m²), detached houses, affluent areas (e.g., Borrowdale, Mount Pleasant). Historically reserved for white residents under colonial-era planning.

Medium Density: Smaller stands (300–1,000 m²), often semi-detached or cluster housing. Mixed-income areas.

High Density: Small stands (<300 m²), often row houses or core houses. Lower-income areas (e.g., Chitungwiza, Glen View). Historically black townships.

The legacy of these density classifications continues to shape property values today, as reflected in our location SHAP values.

K

kVA (Kilovolt-Ampere)

A unit of apparent power, used to rate solar inverters and generators. For example, a "5kVA solar system" can power a typical household. Captured indirectly by `flag_solar`, which matches "kva solar" patterns.

L

Lintel Level

A stage of construction where the walls have reached the height of the window and door lintels (the horizontal supports above openings). A property at lintel level is further along than one at slab level but not yet roofed. Captured in the `construction_stage_ordinal` (score 1 or 2, depending on context).

M

Mabvazuva

A gated community in Ruwa, east of Harare. It is a relatively new, medium-density development. In the case study, House B is located here. The model's location SHAP for Mabvazuva is negative, reflecting its distance from central Harare.

Madokero

A medium-density suburb in Harare West. In the case study, House A is located here. The location SHAP for Madokero is slightly positive relative to the average, reflecting its proximity to amenities.

Mount Pleasant / Mount Pleasant Heights

Affluent, low-density suburbs in Harare North. Mount Pleasant is one of the most expensive areas in the city, with a location SHAP of +0.62. Mount Pleasant Heights is a newer extension with similar characteristics.

P

Prepaid Meter

A electricity meter that requires the user to purchase credit in advance, rather than receiving a monthly bill. Common in Zimbabwe due to cash-flow issues at ZESA. Captured as part of ``flag_zesa_good``.

R

Ruwa

A town east of Harare, administratively separate but functionally part of the greater Harare metropolitan area. It includes suburbs like Mabvazuva, Zimre Park, and Fairview. Property prices in Ruwa are generally lower than in Harare North or West.

S

Shell (Shell Level)

A stage of construction where the structure is complete (walls, roof) but no internal finishes (no ceilings, no kitchen, no paint, bare floors). A shell is less complete than "structurally complete" but more than "slab level." Captured in `construction_stage_ordinal` (score 1).

Slab Level

A stage of construction where only the foundation slab has been laid, and walls may be partially erected. The earliest stage of construction. In the dataset, "slab level" appears in descriptions of incomplete properties. Captured in ``construction_stage_ordinal` (score 1).

Solar (Solar System)

A photovoltaic system that generates electricity from sunlight, often with battery storage. Given Zimbabwe's chronic power shortages, solar is a major selling point. Captured as `flag\_solar`.

Staff Quarters

A separate dwelling on the property for domestic workers, gardeners, or guards. Common in low-density suburbs. Captured as flag_staff_quarters (which also matches "cottage").

Stand

A plot of land, typically with a defined boundary, on which a house is or will be built. In Zimbabwe, "stand size" is synonymous with "land area." The term appears frequently in listings (e.g., "6,000 m² stand"). Captured as `land\_area\_m2`.

T

Tarred Road

Roads surfaced with bitumen or asphalt, as opposed to gravel or dirt. Tarred roads are a sign of higher infrastructure quality and are often mentioned as a selling point. Not explicitly flagged but appears in descriptions (e.g., "tarred roads" in Fairview, Ruwa).

Title Deeds

See "Deeds."

Tynwald

A medium-density suburb in Harare West. In the local SHAP example (HOS240312), Tynwald had a small negative location contribution (-0.02), suggesting it is slightly below the average suburb in terms of price premium.

W

Wall-Plate Level

A stage of construction where the walls have reached the top plate (the horizontal timber or steel beam that supports the roof). This is just before the roof is installed. Captured in construction_stage_ordinal (score 1). The phrase "finish to your own taste" often accompanies wall-plate level properties.

Water Tank

A storage tank, typically plastic or steel, used to store water from a borehole or municipal supply. Essential during water shortages. Captured as flag_water_tank.

Z

ZESA (Zimbabwe Electricity Supply Authority)

The national electricity utility. "ZESA connected" means the property has an active electricity connection. "Good ZESA" implies reliable supply, though this is subjective. Captured as flag_zesa_good. In our model, flag_zesa_good has a small positive effect, as reliable electricity is a valued amenity.

Zimre Park

A suburb in Ruwa, east of Harare. It is a medium-density area with a mix of completed and incomplete properties.

Abbreviations and Acronyms

Abbreviation	Meaning
BIC	Built-In Cupboard
MES	Main Ensuite
ZESA	Zimbabwe Electricity Supply Authority
CBD	Central Business District
CV	Cross-Validation
GLM	Generalised Linear Model
SHAP	SHapley Additive exPlanations
RMSE	Root Mean Squared Error
MAE	Mean Absolute Error
MAPE	Mean Absolute Percentage Error
PI	Prediction Interval
CLIP	Contrastive Language-Image Pre-training
kNN	k-Nearest Neighbours
VIF	Variance Inflation Factor

This glossary is intended to help readers interpret the feature names, model outputs, and case study discussion in the main report. Where a term is directly captured as a feature, the feature name is given in parentheses. Where a term is not explicitly flagged but is relevant to understanding the market context, it is noted as "not explicitly flagged."

The definitions reflect common usage in the Harare property market as observed in the scraped listings. Some terms may have slightly different meanings in other regions of Zimbabwe or in legal contexts. For the purposes of this report, the definitions above apply.